

Distributed Computing and Introduction to High Performance Computing

Imad Kissami¹, Nouredine Ouhaddou¹

¹Mohammed VI Polytechnic University, Benguerir, Morocco



Outline of this lecture

- About Python
- Python is slow
- Profiling a Python code

About Python

- Python was created by Guido van Rossum in 1991 (last version 3.9 - 05/10/2020)
- Python is **simple**
- Python is **fully featured**
- Python is **readable**
- Python is **extensible**
- Python is **ubiquitous, portable, and free**
- Python has **many third party libraries, tools, and a large community**

About Python

- Python was created by Guido van Rossum in 1991 (last version 3.9 - 05/10/2020)
 - Python is **simple**
 - Python is **fully featured**
 - Python is **readable**
 - Python is **extensible**
 - Python is **ubiquitous, portable, and free**
 - Python has **many third party libraries, tools, and a large community**
- ➡ **But Python is slow!!**

About Python

- Python was created by Guido van Rossum in 1991 (last version 3.9 - 05/10/2020)
 - Python is **simple**
 - Python is **fully featured**
 - Python is **readable**
 - Python is **extensible**
 - Python is **ubiquitous, portable, and free**
 - Python has **many third party libraries, tools, and a large community**
- ➡ **Does is really matters?**

Python is slow

When does it matter?

- Is my code fast enough to produce the results I need in the time I have?
- How many CPUh is this code going to waste over its lifetime?
 - How inefficient is it?
 - How long does it run?
 - How often will it run?
- Does it cause problems on the system it's running on?
- How much effort is it to make it run faster?
- ➡ For those who are interested, you can follow this MOOC

Profiling a Python code

Why?

- It's mandatory to know what sections of code are bottlenecks in order to improve performance.
- You need to measure it, not to guess it
- **Premature optimization is the root of all evil** D. Knuth
- **First make it work. Then make it right. Then make it fast.** K. Beck
- How?

Profiling a Python code

Different kinds of profilers

- Deterministic and statistical profiling
 - the profiler will be monitoring all the events
 - it will sample after time intervals to collect that information
- The level at which resources are measured; module, function or line level
- Profile visualizers

Profiling a Python code

Available tools

- **Inbuilt timing modules**
- **profile and cProfile**
- pstats
- **line_profiler**
- Yappi
- vmprof-python
- pyinstrument
- gprof2dot
- pyprof2calltree, KCacheGrind
- **snakeviz**
- Scalene

Profiling a Python code

Use case

```
1 def linspace(start, stop, n):
2     step = float(stop - start) / (n - 1)
3     return [start + i * step for i in range(n)]
4
5 def mandel(c, maxiter):
6     z = c
7     for n in range(maxiter):
8         if abs(z) > 2:
9             return n
10        z = z*z + c
11    return n
12
13 def mandel_set(xmin=-2.0, xmax=0.5, ymin=-1.25, ymax=1.25,
14               width=1000, height=1000, maxiter=80):
15     r = linspace(xmin, xmax, width)
16     i = linspace(ymin, ymax, height)
17     n = [[0]*width for _ in range(height)]
18     for x in range(width):
19         for y in range(height):
20             n[y][x] = mandel(complex(r[x], i[y]), maxiter)
21     return n
```

Profiling a Python code

timeit

The very naive way

```
1 import time
2
3 start_time = time.time()
4 mandel_set()
5 end_time = time.time()
6 # Time taken in seconds
7 elapsed_time = end_time - start_time
8
9 print('> Elapsed time', elapsed_time)
```

or using the magic method timeit

```
1 [In] %timeit mandel_set()
2 [Out] 3.01 s +/- 84.6 ms per loop (mean +/- std. dev. of 7 runs, 1 loop each)
```

Profiling a Python code

prun

```
1 [In] %prun -s cumulative mandel_set()
```

which is, in console mode, equivalent to

```
1 python -m cProfile -s cumulative mandel.py
```

```
1          25214601 function calls in 5.151 seconds
2
3      Ordered by: cumulative time
4
5      ncalls tottime percall cumtime percall filename:lineno(function)
6          1 0.000 0.000 5.151 5.151 {built-in method builtins.exec}
7          1 0.002 0.002 5.151 5.151 <string>:1(<module>)
8          1 0.291 0.291 5.149 5.149 <ipython-input-4-9421bc2016cb>:13(mandel_set)
9      1000000 3.461 0.000 4.849 0.000 <ipython-input-4-9421bc2016cb>:5(mandel)
10 24214592 1.388 0.000 1.388 0.000 {built-in method builtins.abs}
11          1 0.008 0.008 0.008 0.008 <ipython-input-4-9421bc2016cb>:17(<listcomp>)
12          2 0.000 0.000 0.000 0.000 <ipython-input-4-9421bc2016cb>:1(linspace)
13          2 0.000 0.000 0.000 0.000 <ipython-input-4-9421bc2016cb>:3(<listcomp>)
14          1 0.000 0.000 0.000 0.000 {method 'disable' of '_lsprof.Profiler' objects}
```

- Most of the time is spent in the mandel function
- profiling introduces some overhead 5.14(s) instead of 3.01(s)

Profiling a Python code

Visualization

- Profiling results can be visualized with SnakeViz
- We must be in console mode

```
1 python3 -m cProfile -o mandel.prof mandel.py
2 snakeviz --port 6542 --hostname localhost --server mandel.prof
```

SnakeViz

Reset Root

Reset Zoom

Style: Icicle

Depth: 10

Coeff: 1 / 1000

Name:

mandel

Cumulative Time:


1.75 s (100.00 %)

File:

mandel.py

Lines:

Directory:



scalls	tottime	percall	cumtime	percall	filename:line(function)
1050000	3.415	3.415e-06	4.807	4.807e-06	mandel.py:9(mandel)
24214688	1.393	5.752e-08	1.393	5.752e-08	--(<built-in method builtins.abc>)
1	0.2763	0.2763	5.087	5.087	mandel.py:17(mandel_set)
316	0.03232	0.0001023	0.03232	0.0001023	--(<built-in method builtins.compile>)

Profiling a Python code

Details at the line level

- We know that most of the time is spent in the `mandel` function
- We shall use the `line_profiler` package on this function to get details at the line level

```
1 [In] %load_ext line_profiler
2 [In] %lprun -f mandel mandel_set()
```

Which gives the result

Timer unit: 1e-06 s

Total time: 22.8401 s

File: <ipython-input-2-9421bc2016cb>

Function: `mandel` at line 5

#Line	Hits	Time	Per Hit	% Time	Line Contents
5					<code>def mandel(c, maxiter):</code>
6	1000000	250304.0	0.3	1.1	<code>z = c</code>
7	24463110	6337732.0	0.3	27.7	<code>for n in range(maxiter):</code>
8	24214592	8327289.0	0.3	36.5	<code>if abs(z) > 2:</code>
9	751482	201108.0	0.3	0.9	<code>return n</code>
10	23463110	7658255.0	0.3	33.5	<code>z = z*z + c</code>
11	248518	65444.0	0.3	0.3	<code>return n</code>

Profiling a Python code

Details at the line level

This can be done in console mode as well

```
1 import line_profiler
2
3 @profile
4 def mandel(c, maxiter):
5     z = c
6     for n in range(maxiter):
7         if abs(z) > 2:
8             return n
9         z = z*z + c
10    return n
```

Then on the command line

```
1 kernprof -l -v mandel.py
```