

Distributed Computing and Introduction to High Performance Computing

Ahmed Ratnani¹

¹Mohammed VI Polytechnic University, Benguerir, Morocco



Outline of this lecture

- Some basics on Programming Languages and Compilers
- What is Numba?
- What is Pyccel?
- ➡ *see the last Pyccel talk*

A brief history on programming languages

- Lambda Calculus (A. Church - 1930s): formal system for expressing computation based on functions
- Turing machine (1936): Abstract machine for computing (abstract) functions
- Von Neumann model (1945): first model for computers
- First high-level language:
 - Fortran (J. Backus - 1954)
 - Lisp (J. Mc Carthy - 1958)
- Two paradigms emerged:
 - imperative programming
 - functional programming

Language Theory

A brief introduction to Compilers

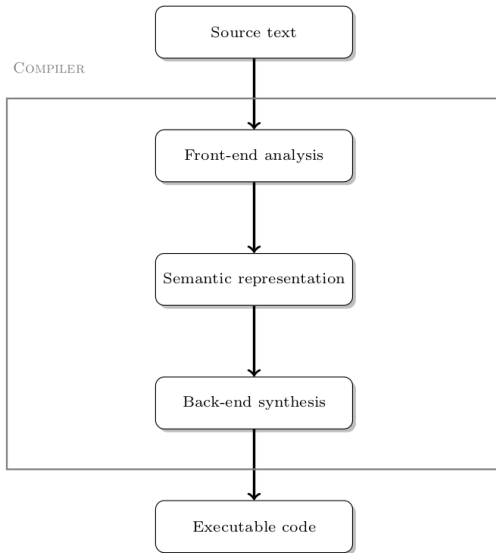


Figure 1: Conceptual structure of a compiler

Language Theory

A brief introduction to Compilers

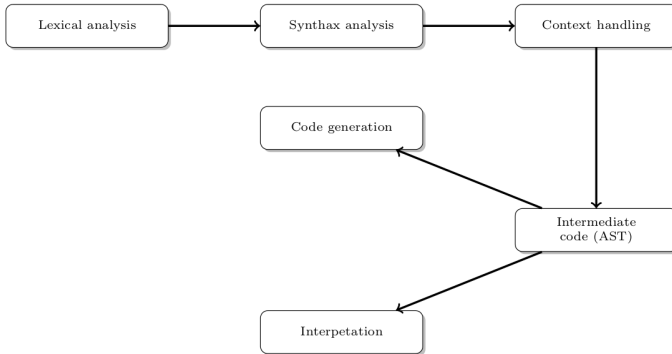


Figure 2: Conceptual structure of a compiler

Numba

What is Numba?

The translation/magic is been done using the LLVM compiler, which is open sourced and has quite active dev community.

- Numba compiles Python code at runtime to native machine instructions
- Unlike Cython, it is backward compatible (same for Pyccel and Pythran)
- It is based on the LLVM compiler
- Open source code with a big active community

Installation

```
1 pip install numba
```

Typical usage

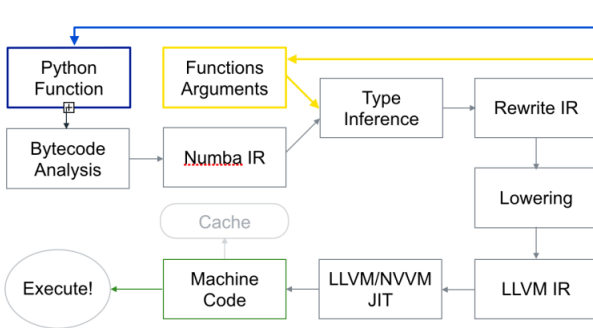
```
1 @jit
2 def myfunction(*args, **kwargs):
3     ...
```



Numba

Workflow

```
@jit
def do_math(a, b):
    ...
>>> do_math(x, y)
```



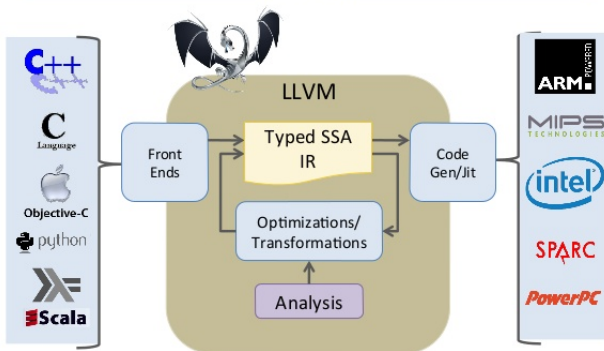
Numba

Current limitations

- Numba compiles Python functions, not entire programs (pypy is great for that). It also doesn't support partial compilation of functions – it needs to be able to resolve all data types in the selected function.
- Presently, Numba is focused on numerical data types, like int, float, and complex. There is very limited string processing support and the best results are realised with Numpy arrays.
- Decorating functions that make use of Pandas (or other unsupported data structures) would deteriorate performance. Pandas is not understood by Numba and as a result, Numba would simply run this code via the interpreter but with the additional cost of the Numba internal overheads.
- You are better off using Cython for code that interferes with C++, as Numba can't talk with C++ effectively unless a C wrapper is used.
- Numba doesn't generate C/C++ code that can be used for a separate compilation; it goes directly from Python down to LLVM code. Cython would be more suitable for this use case, as it allows inspection of the code in C++ before compilation.

LLVM Compiler Infrastructure

[Lattner et al.]



■ The LLVM Compiler Infrastructure

- Provides reusable components for building compilers
- Reduce the time/cost to build a new compiler
- Build static compilers, JITs, trace-based optimizers, ...

■ The LLVM Compiler Framework

- End-to-end compilers using the LLVM infrastructure
- C and C++ gcc frontend
- Backends for C, X86, Sparc, PowerPC, Alpha, Arm, Thumb, IA-64, ...

- The LLVM Virtual Instruction Set
 - The common language- and target-independent IR
 - Internal (IR) and external (persistent) representation
- A collection of well-integrated libraries
 - Analyses, optimizations, code generators, JIT compiler, garbage collection support, profiling, ...
- A collection of tools built from the libraries
 - Assemblers, automatic debugger, linker, code generator, compiler driver, modular optimizer, ...

- Analyze and optimize as early as possible
 - Compile-time opts reduce modify-rebuild-execute cycle
 - Compile-time optimizations reduce work at link-time (by shrinking the program)
- One IR (without lowering) for analysis and optimization
 - Compile-time optimizations can be run at link-time too
 - The same IR is used as input to the JIT

- Easy to produce, understand, and define
- Language- and Target-Independent
 - AST-level IR is not very feasible
 - Every analysis/transformation must know about all languages
- One IR for analysis and optimization
 - R must be able to support aggressive IPO, loop opts, scalar opts, ... high- and low-level optimization
- Optimize as much as early as possible
 - Can't postpone everything until link or runtime
 - No lowering in the IR